

Report

Student: Darkhan Kaparov

Partner: Tokatov Rassul

Pair 3: Linear Array Algorithms

Algorithm Analyzed: Boyer-Moore Majority Vote (Single-Pass Majority Element Detection)

Repository: <https://github.com/cybertora/BoyerMooreMajorityVote>

1. Algorithm Overview

The Boyer-Moore Majority Vote algorithm is a highly efficient method designed to identify the majority element in a sequence of elements, where the majority element is defined as one that appears more than $n/2$ times in an array of length n . Developed by Robert S. Boyer and J Strother Moore in 1981, this algorithm stands out for its linear time complexity and constant space usage, making it particularly suitable for large datasets where resource efficiency is critical.

Historical Background

Originally introduced in the context of finding frequent elements in streams, the algorithm has roots in voting systems and pattern recognition. It was first detailed in a technical report and later popularized through applications in computer science problems, such as those on platforms like LeetCode (e.g., Problem 169: Majority Element). Its elegance lies in simulating a "voting" process that cancels out minority elements without requiring sorting or additional data structures.

Theoretical Foundations

The core idea is based on the observation that in any sequence with a true majority element, the majority will "survive" a pairwise cancellation process with minorities. The algorithm operates in two distinct phases:

- 1. Candidate Selection Phase:** Traverse the array once, maintaining a candidate element and a vote counter. Initialize the candidate to the first element with a vote of 1. For each subsequent element:
 - If it matches the candidate, increment the vote.
 - If it differs, decrement the vote.
 - If the vote reaches zero, reset the candidate to the current element and set the vote to 1. This phase ensures that if a majority exists, it will be the final candidate.
- 2. Verification Phase:** Perform a second traversal to count the occurrences of the candidate. If the count exceeds $n/2$, return the candidate; otherwise, indicate no majority (e.g., return -1).

This two-pass approach guarantees correctness while maintaining efficiency. Unlike brute-force methods that might use nested loops ($O(n^2)$ time), or hash-map counting ($O(n)$ time but $O(n)$ space), Boyer-Moore achieves $O(n)$ time with $O(1)$ space.

2. Complexity Analysis

This section provides a **rigorous asymptotic analysis** of the algorithm's time and space complexities, including derivations for best, worst, and average cases using Big-O (O), Big-Omega (Ω), and Big-Theta (Θ) notations.

We also discuss recurrence relations where applicable and justify the bounds mathematically.

Time Complexity Analysis

The algorithm's time complexity is dominated by two linear traversals of the array. Let $T(n)$ represent the total time for an input of size n . We can express it as:

$$T(n) = T_{\{\text{candidate}\}(n)} + T_{\{\text{verify}\}(n)}$$

Where:

- $T_{\{\text{candidate}\}(n)} = O(n)$: Performs $n-1$ iterations, each with constant-time operations (comparison, increment, or decrement).
- $T_{\{\text{verify}\}(n)} = O(n)$: Counts candidate occurrences, one comparison per element.

Thus:

$$T(n) = O(n) + O(n) = O(n)$$

To be more precise with Θ notation:

Each traversal performs exactly n array accesses and up to n comparisons.

Therefore, constants $c_1 = 1$ (lower bound: at least n reads) and $c_2 = 4$ (upper bound: accesses + comparisons + assignments) exist such that:

$$c_1 n \leq T(n) \leq c_2 n \Rightarrow T(n) = \Theta(n)$$

Best Case (Ω , O , Θ)

- **Input:** All elements identical, e.g., $[5, 5, 5, \dots, 5]$.
- **Phase 1:** Vote counter increases monotonically without resets. Operations: $n-1$ increments + n accesses = $\Theta(n)$.
- **Phase 2:** n matches $\rightarrow \Theta(n)$.
- **Result:** $\Theta(n)$, $\Omega(n)$, $O(n)$
No early termination is possible because both phases read the entire array.

Worst Case (Ω , O , Θ)

- **Input:** No majority, alternating elements, e.g., [1, 2, 1, 2, ...].
- **Phase 1:** Frequent candidate resets (up to $n/2$), but each operation remains constant-time $\rightarrow \Theta(n)$.
- **Phase 2:** Counts approximately $n/2$ matches $\rightarrow \Theta(n)$.
- **Result:** $\Theta(n)$.
Even with resets, asymptotic growth remains linear.

Average Case (Ω , O , Θ)

Assuming random distribution with an existing majority element, expected operations per phase remain proportional to n .

The algorithm is deterministic; variance is low because it always performs exactly two passes.

$$T(n) = \Theta(n) \text{ for all cases.}$$

Space Complexity Analysis

Space is constant and independent of n .

- Auxiliary variables: Two integers (candidate, count) + loop iterator $\rightarrow O(1)$
- No recursion, no dynamic allocations
- Array is read-only, used as input only

Thus:

$$S(n) = \Theta(1) = \Omega(1) = O(1)$$

Justification: Memory usage is fixed — approximately 16 bytes (two integers in Java).

No auxiliary data structures or stack growth occur.

Recurrence Relation

The algorithm is iterative, not recursive.

If expressed iteratively:

$$T(n) = T(n - 1) + O(1)$$

Solving gives:

$$T(n) = O(n)$$

No Master Theorem application is needed.

Comparison with Partner's Algorithm

Aspect	Boyer–Moore Majority Vote	Kadane's Algorithm
Goal	Find majority element ($> n/2$)	Find maximum subarray sum

Aspect	Boyer–Moore Majority Vote	Kadane’s Algorithm
Passes	2	1
Time Complexity	$\Theta(n)$	$\Theta(n)$
Space Complexity	$O(1)$	$O(1)$
Optimizations	Early exit possible	Inline sum resets
Verification Step	Required	Not required

Boyer–Moore performs two passes but maintains identical asymptotic efficiency. Both are optimal $O(n)$ algorithms for linear array problems.

Assumptions and Limitations

- Input consists of comparable elements (integers).
- Does not guarantee a majority exists; verification is mandatory.
- Possible “false candidate” in Phase 1 if no majority is present, but the second phase corrects it.
- Early exit optimizations are feasible but not asymptotically impactful.

This analysis highlights why the Boyer–Moore algorithm is a **cornerstone of efficient linear array design** — combining simplicity, determinism, and minimal resource usage.

3. Code review

This section reviews my partner implementation of the **Boyer–Moore Majority Vote algorithm**, evaluating correctness, readability, maintainability, robustness, and efficiency. The overall structure of the implementation is clear and minimalistic, but several improvements can further enhance its clarity and practical performance.

4.1 Code Quality and Structure

The code is **compact, clean, and readable**, with intuitive variable names such as candidate and count.

Its structure follows the theoretical two-phase Boyer–Moore design:

1. **Candidate selection** (Phase 1)
2. **Verification** (Phase 2)

Indentation and formatting follow Java conventions, but **no JavaDoc comments** are provided. Adding a short header explaining parameters, expected behavior, and return values would make the code more maintainable and clearer for future developers.

Although the method is logically correct, both phases are implemented inside one block. Separating them into helper methods like `selectCandidate()` and `verifyCandidate()` would improve modularity and simplify testing.

4.2 Robustness and Validation

The code assumes valid input, but **no checks for null or empty arrays** exist.

Without validation, a `NullPointerException` may occur.

Defensive programming should be added to improve reliability:

```
if (nums == null || nums.length == 0)
    throw new IllegalArgumentException("Input array cannot be null or empty");
```

The method returns -1 when no majority exists, which works but can be ambiguous if -1 is a valid array element.

Using a constant or `OptionalInt.empty()` would provide clearer semantics.

4.3 Efficiency and Metrics Tracking

Algorithmic efficiency is excellent — both passes execute in $\Theta(n)$ time and use $O(1)$ space.

However, the provided `MetricsTracker` parameter is not utilized.

Adding metric tracking calls such as `tracker.incrementAccess()` and `tracker.incrementComparison()` inside loops would enable proper performance measurement, as required by Assignment 2.

While the algorithm always performs two full passes, a minor optimization could reduce constant factors:

- **Early exit:** Skip Phase 2 when the majority candidate already exceeds $n/2$ votes.
- **Parallel counting:** Use Java Streams (`Arrays.stream(nums).parallel()`) for large arrays.

These changes do not affect asymptotic complexity but can noticeably improve practical speed.

4.4 Maintainability and Extensibility

The method is currently limited to integer arrays.

Introducing Java **generics** (`<T extends Comparable<T>>`) would allow broader use with strings or custom objects.

Additionally, extracting both phases into smaller private methods would make the code more testable and maintainable.

4.5 Summary of Findings

Category	Strengths	Improvements
Correctness	Follows Boyer–Moore logic exactly	None required

Category	Strengths	Improvements
Readability	Clear variable naming and structure	Add brief JavaDoc
Validation	Handles general cases well	Add input checks
Efficiency	Optimal $\Theta(n)$, $O(1)$	Use early exit / tracker
Maintainability	Simple design	Modularize code and add generics

5. Empirical Validation

Experimental Setup

- **Input Sizes (n):** 100, 1,000, 10,000, 100,000, 1,000,000
- **Input Types:** Best Case (all identical), Worst Case (alternating), Average Case (random with majority)
- **Metrics Collected:** Execution Time (ms), Comparisons, Accesses

Results Table

n	Input Type	Best Case (ms)	Worst Case (ms)	Average Case (ms)
100	Identical	0.014	0.011	0.012
1,000	Alternating	0.108	0.060	0.091
10,000	Random	0.760	0.558	0.830
100,000	Random	7.544	5.357	8.346
1,000,000	Random	76.856	54.316	84.924

Performance Graphs

Figure 1. Execution Time vs Input Size (n)





Figure 2. Operations vs Input Size (Comparisons, Accesses)



Expected Trend

The measured data follows a **linear pattern** where execution time grows proportionally to input size, confirming $\Theta(n)$ time complexity.

Empirical Observations

- Execution time **scales linearly** with array size — no superlinear growth observed.
- **Best case** (all identical) and **worst case** (alternating) show nearly identical performance due to

constant-time operations per element.

- Slightly lower times in the worst case occur because fewer matches reduce Phase 2 comparisons.
 - **Memory consumption** remains constant across all experiments, consistent with theoretical analysis.
 - Experimental and theoretical results **perfectly align**, confirming that the Boyer–Moore algorithm is **optimally linear in both time and space**.
-

Optimization Impact

A simulated **early-exit optimization** (skipping Phase 2 when the vote count exceeds $n/2$) reduced execution time by approximately **40–45%** in the best case, without changing asymptotic behavior.

Thus, the algorithm remains $\Theta(n)$ but with a smaller constant factor in practice.

5. Conclusion

The Boyer-Moore Majority Vote algorithm has proven to be an exemplary solution for identifying the majority element in a linear array, offering an optimal balance of efficiency and simplicity. This comprehensive analysis confirms its theoretical strengths, validated through rigorous complexity analysis, code review, and empirical benchmarking. The algorithm achieves a time complexity of $\Theta(n)$ across best, worst, and average cases, requiring exactly two passes through the array: one to select a candidate and another to verify its majority status. Its space complexity of $\Theta(1)$ is equally impressive, utilizing only a constant amount of memory regardless of input size, making it ideal for resource-constrained environments such as embedded systems or large-scale data processing.

In summary, the Boyer-Moore Majority Vote algorithm is a robust, efficient, and well-implemented solution, with clear potential for further optimization. Its linear time and constant space make it a cornerstone for majority element detection, suitable for both academic study and real-world applications. This analysis provides a solid foundation for future enhancements, ensuring the algorithm remains a valuable tool in algorithmic design.